

Airflow operators

INTRODUCTION TO APACHE AIRFLOW IN PYTHON



Mike Metzger
Data Engineer

Operators

- Represent a single task in a workflow.
- Run independently (usually).
- Generally do not share information.
- Various operators to perform different tasks.

```
# New way, Airflow 2.x+  
EmptyOperator(task_id='example')  
  
# Old way, Airflow <2.0  
EmptyOperator(task_id='example', dag=dag_name)
```

BashOperator

```
BashOperator(  
    task_id='bash_example',  
    bash_command='echo "Example!"',  
    # Next line only for Airflow before v2.0  
    dag=dag  
)
```

```
BashOperator(  
    task_id='bash_script_example',  
    bash_command='runcleanup.sh',  
)
```

- Executes a given Bash command or script.
- Runs the command in a temporary directory.
- Can specify environment variables for the command.

BashOperator examples

```
from airflow.operators.bash import BashOperator
example_task = BashOperator(task_id='bash_ex',
                             bash_command='echo 1',
                             )
```

```
bash_task = BashOperator(task_id='clean_addresses',
                           bash_command='cat addresses.txt | awk "NF==10" > cleaned.txt',
                           )
```

Operator gotchas

- Not guaranteed to run in the same location / environment.
- May require extensive use of Environment variables.
- Can be difficult to run tasks with elevated privileges.

Let's practice!

INTRODUCTION TO APACHE AIRFLOW IN PYTHON

Airflow tasks

INTRODUCTION TO APACHE AIRFLOW IN PYTHON



Mike Metzger
Data Engineer

Tasks

Tasks are:

- Instances of operators
- Usually assigned to a variable in Python

```
example_task = BashOperator(task_id='bash_example',  
                             bash_command='echo "Example!"')
```

- Referred to by the task_id within the Airflow tools

Task dependencies

- Define a given order of task completion
- Are not required for a given workflow, but usually present in most
- Are referred to as *upstream* or *downstream* tasks
- In Airflow 1.8 and later, are defined using the *bitshift* operators
 - `>>`, or the upstream operator
 - `<<`, or the downstream operator

Upstream vs Downstream

Upstream means **before**

Downstream means **after**


Simple task dependency

```
# Define the tasks
task1 = BashOperator(task_id='first_task',
                      bash_command='echo 1'
                      )

task2 = BashOperator(task_id='second_task',
                      bash_command='echo 2'
                      )

# Set first_task to run before second_task
task1 >> task2    # or task2 << task1
```

Task dependencies in the Airflow UI

 Airflow

DAGsCluster ActivityDatasetsSecurity▼Browse▼Admin▼Docs▼01:41 UTC ▼→] Log In

☐ DAG: simple_dependency

Schedule: 1 day, 0:00:00Next Run: 2024-01-10, 00:00:00

Grid

Graph

Calendar

Task Duration

Task Tries

Landing Times

Gantt

Details

<> Code

Audit Log

01/30/2024, 01:41:01 AM

25

All Run Types

All Run States

Clear Filters

Auto-refresh

Press **shift** + **/** for Shortcuts

deferred

failed

queued

removed

restarting

running

scheduled

shutdown

skipped

success

up_for_reschedule

up_for_retry

upstream_failed

no_status

« » DAG simple_dependency

Details

Graph

Gantt

<> Code

first_task

BashOperator

second_task

BashOperator


Layout:

Top -> Bottom

first_task

second_task

Task dependencies in the Airflow UI

 Airflow

DAGsCluster ActivityDatasetsSecurity▼Browse▼Admin▼Docs▼01:41 UTC → Log In

☐ DAG: simple_dependency

Schedule: 1 day, 0:00:00Next Run: 2024-01-10, 00:00:00

Grid

Graph

Calendar

Task Duration

Task Tries

Landing Times

Gantt

Details

Code

Audit Log

01/30/2024, 01:41:01 AM

25

All Run Types

All Run States

Clear Filters

Auto-refresh

Press **shift** + **/** for Shortcuts

deferred

failed

queued

removed

restarting

running

scheduled

shutdown

skipped

success

up_for_reschedule

up_for_retry

upstream_failed

no_status

« » DAG simple_dependency

Details

Graph

Gantt

Code

first_task
BashOperator


second_task
BashOperator

Layout:
Top -> Bottom

first_task

second_task

Task dependencies in the Airflow UI

 Airflow

DAGsCluster ActivityDatasetsSecurity▼Browse▼Admin▼Docs▼01:43 UTC → Log In

☐ DAG: simple_dependency

Schedule: 1 day, 0:00:00Next Run: 2024-01-10, 00:00:00

Grid

Graph

Calendar

Task Duration

Task Tries

Landing Times

Gantt

Details

Code

Audit Log

01/30/2024, 01:42:44 AM

25

All Run Types

All Run States

Clear Filters

Auto-refresh

Press **shift** + **/** for Shortcuts

deferredfailedqueuedremovedrestartingrunningscheduledshutdownskippedsuccessup_for_rescheduleup_for_retryupstream_failedno_status

« » DAGsimple_dependency

Details

Graph

Gantt

Code

first_task
BashOperator

second_task
BashOperator

Layout:
Left -> Right

first_task

second_task

Multiple dependencies

Chained dependencies:

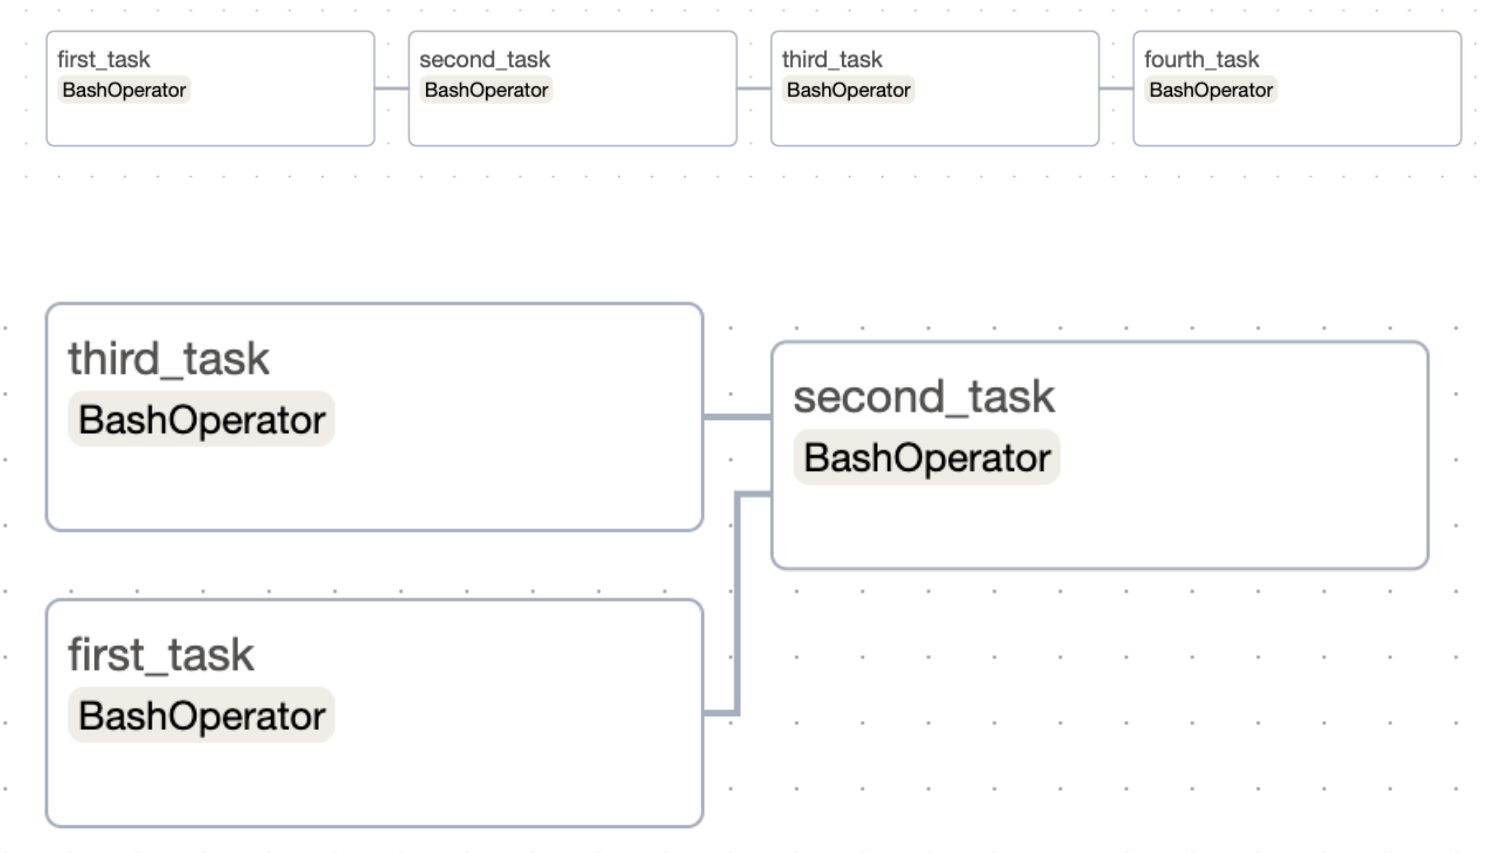
```
task1 >> task2 >> task3 >> task4
```

Mixed dependencies:

```
task1 >> task2 << task3
```

or:

```
task1 >> task2  
task3 >> task2
```



Let's practice!

INTRODUCTION TO APACHE AIRFLOW IN PYTHON

Additional operators

INTRODUCTION TO APACHE AIRFLOW IN PYTHON



Mike Metzger
Data Engineer

PythonOperator

- Executes a Python function / callable
- Operates similarly to the BashOperator, with more options
- Can pass in arguments to the Python code

```
from airflow.operators.python import PythonOperator
def printme():
    print("This goes in the logs!")
python_task = PythonOperator(
    task_id='simple_print',
    python_callable=printme
)
```

Arguments

- Supports arguments to tasks
 - Positional
 - Keyword
- Use the `op_kwargs` dictionary

op_kwargs example

```
def sleep(length_of_time):  
    time.sleep(length_of_time)  
  
sleep_task = PythonOperator(  
    task_id='sleep',  
    python_callable=sleep,  
    op_kwargs={'length_of_time': 5}  
)
```

EmailOperator

- Found in the `airflow.operators` library
- Sends an email
- Can contain typical components
 - HTML content
 - Attachments
- Does require the Airflow system to be configured with email server details

EmailOperator example

```
from airflow.operators.email import EmailOperator

email_task = EmailOperator(
    task_id='email_sales_report',
    to='sales_manager@example.com',
    subject='Automated Sales Report',
    html_content='Attached is the latest sales report',
    files='latest_sales.xlsx'
)
```

Let's practice!

INTRODUCTION TO APACHE AIRFLOW IN PYTHON

Airflow scheduling

INTRODUCTION TO APACHE AIRFLOW IN PYTHON




Mike Metzger
Data Engineer

DAG Runs

- A specific instance of a workflow at a point in time
- Can be run manually or via `schedule_interval`
- Maintain state for each workflow and the tasks within
 - `running`
 - `failed`
 - `success`

¹ <https://airflow.apache.org/docs/stable/scheduler.html>

DAG Runs view

 Airflow

DAGs

Cluster Activity

Datasets

Security

Browse

DAG Runs

Jobs

Audit Logs

Task Instances

Task Reschedules

Triggers

SLA Misses

DAG Dependencies

Admin

Docs

04:39 UTC

→ Log In

List Dag Run

Search


Actions

←

Record Count: 5

<input type="checkbox"/>	State	Dag Id	Logical Date	Run Id	State	Queued At	Start Date
<input type="checkbox"/>	running	example_dag	2024-01-12, 00:00:00	scheduled__2024-01-12T00:00:00+00:00	scheduled	2024-01-29, 04:38:41	2024-01-29, 04:38:41
<input type="checkbox"/>	success	example_dag	2024-01-11, 00:00:00	scheduled__2024-01-11T00:00:00+00:00	scheduled	2024-01-29, 04:38:30	2024-01-29, 04:38:30
<input type="checkbox"/>	failed	update_state	2024-01-11, 00:00:00	scheduled__2024-01-11T00:00:00+00:00	scheduled	2024-01-29, 04:38:36	2024-01-29, 04:38:36

DAG Runs state

 Airflow

DAGs

Cluster Activity

Datasets

Security

Browse

Admin

Docs

04:39 UTC

Log In

List Dag Run

Search

Actions

Record Count: 5

	State	Dag Id	Logical Date	Run Id	Run Type	Queued At	Start Date
<input type="checkbox"/>	running	example_dag	2024-01-12, 00:00:00	scheduled__2024-01-12T00:00:00+00:00	scheduled	2024-01-29, 04:38:41	2024-01-29, 04:38:41
<input type="checkbox"/>	success	example_dag	2024-01-11, 00:00:00	scheduled__2024-01-11T00:00:00+00:00	scheduled	2024-01-29, 04:38:30	2024-01-29, 04:38:30
<input type="checkbox"/>	failed	update_state	2024-01-11, 00:00:00	scheduled__2024-01-11T00:00:00+00:00	scheduled	2024-01-29, 04:38:36	2024-01-29, 04:38:36

Schedule details

When scheduling a DAG, there are several attributes of note:

- `start_date` - The date / time to initially schedule the DAG run
- `end_date` - Optional attribute for when to stop running new DAG instances
- `max_tries` - Optional attribute for how many attempts to make
- `schedule_interval` - How often to run

Schedule interval

`schedule_interval` represents:

- How often to schedule the DAG
- Between the `start_date` and `end_date`
- Can be defined via `cron` style syntax or via built-in presets.

cron syntax

```
# |----- minute (0 - 59)
# | |----- hour (0 - 23)
# | | |----- day of the month (1 - 31)
# | | | |----- month (1 - 12)
# | | | | |----- day of the week (0 - 6) (Sunday to Saturday;
# | | | | |                                     7 is also Sunday on some systems)
# | | | | |
# | | | | |
# * * * * * command to execute
```

- Is pulled from the Unix cron format
- Consists of 5 fields separated by a space
- An asterisk `*` represents running for every interval (ie, every minute, every day, etc)
- Can be comma separated values in fields for a list of values

cron examples

```
0 12 * * * # Run daily at noon
```

```
* * 25 2 * # Run once per minute on February 25
```

```
0,15,30,45 * * * * # Run every 15 minutes
```

Airflow scheduler presets

Preset:

- @hourly
- @daily
- @weekly
- @monthly
- @yearly

cron equivalent:

- 0 * * * *
- 0 0 * * *
- 0 0 * * 0
- 0 0 1 * *
- 0 0 1 1 *

¹ <https://airflow.apache.org/docs/stable/scheduler.html>

Special presets

Airflow has two special `schedule_interval` presets:

- `None` - Don't schedule ever, used for manually triggered DAGs
- `@once` - Schedule only once

schedule_interval issues

When scheduling a DAG, Airflow will:

- Use the `start_date` as the earliest possible value
- Schedule the task at `start_date` + `schedule_interval`

```
'start_date': datetime(2020, 2, 25),  
'schedule_interval': @daily
```

This means the earliest starting time to run the DAG is on February 26th, 2020

Let's practice!

INTRODUCTION TO APACHE AIRFLOW IN PYTHON